

Handwritten Letter Prediction

3/1/2022
Cory Randolph

A dark blue diagonal gradient bar that starts from the bottom left corner and extends towards the top right corner, covering the lower half of the slide.


Data Overview

- Imported EMNIST- balanced dataset.
- The EMNIST dataset has a set of handwritten digits (0–9), (a-z), and (A-Z) being converted into 28x28 pixel pictures.
- The EMNIST Balanced dataset - 47-class dataset was chosen over the By Class dataset to avoid classification errors.
- This project disregarded the numerical (0-9) classes to focus just on letters


Label: a



Label: n



Label: F



Modeling

1. **Building the model-** Using Keras- high level API of tensorflow framework build a layered sequential type

Model.- Based on **Convolutional Neural Network architecture**. model type= Sequential()

```
model.add(Conv2D(32, kernel_size=3, activation='relu', input_shape=(28, 28, 1)))
```

2. **Compiling the model- parameters [optimizer, Loss, Metrics]**

```
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

```
model = Sequential()
```

```
model.add(Conv2D(32, kernel_size=3, activation='relu', input_shape=(28, 28, 1)))  
model.add(BatchNormalization())  
model.add(Conv2D(32, kernel_size=3, activation='relu'))  
model.add(BatchNormalization())  
model.add(Conv2D(32, kernel_size=5, strides=2, padding='same', activation='relu'))  
model.add(BatchNormalization())  
model.add(Dropout(0.4))
```


3. **Training the Model**

```
history = model.fit(x = x_train,  
                    y = y_train,  
                    validation_data = (x_test, y_test),  
                    epochs=20) #10
```

4. **Evaluation and Prediction**

```
model.evaluate(x_test, y_test)
```

Predicted Label: g
Actual Label: g



Model Architecture

```
model = Sequential()

model.add(Conv2D(32, kernel_size=3, activation='relu', input_shape=(28, 28, 1)))
model.add(BatchNormalization())
model.add(Conv2D(32, kernel_size=3, activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, kernel_size=5, strides=2, padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))

model.add(Conv2D(64, kernel_size=3, activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, kernel_size=3, activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, kernel_size=5, strides=2, padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.4))
model.add(Dense(37, activation='softmax')) #26 fro just Capital Letters
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
batch_normalization (Batch Normalization)	(None, 26, 26, 32)	128
conv2d_1 (Conv2D)	(None, 24, 24, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 24, 24, 32)	128
conv2d_2 (Conv2D)	(None, 12, 12, 32)	25632
batch_normalization_2 (Batch Normalization)	(None, 12, 12, 32)	128
dropout (Dropout)	(None, 12, 12, 32)	0
conv2d_3 (Conv2D)	(None, 10, 10, 64)	18496
batch_normalization_3 (Batch Normalization)	(None, 10, 10, 64)	256
conv2d_4 (Conv2D)	(None, 8, 8, 64)	36928
batch_normalization_4 (Batch Normalization)	(None, 8, 8, 64)	256
conv2d_5 (Conv2D)	(None, 4, 4, 64)	102464
batch_normalization_5 (Batch Normalization)	(None, 4, 4, 64)	256
dropout_1 (Dropout)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
batch_normalization_6 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 37)	4773
Total params: 330,725		
Trainable params: 329,893		
Non-trainable params: 832		


Evaluation

Evaluated on test data set. Using below measures:


- Confusion Matrix
- Accuracy
- Loss
- Precision
- Recall
- f1-score

```
[ ] cm = confusion_matrix(y_test_as_labels, y_test_preds_as_labels)
```

```
plt.figure(figsize=[10,10])  
sns.heatmap(cm, cmap="Reds", annot=True, fmt='.0f', xticklabels = labels_dict_caps.values(), yticklabels= labels_dict_caps.values())  
plt.show()
```



Evaluation



Front End Design (Streamlit)

Main Features

- User draws a letter
- Custom ML model makes a class/letter prediction
- Probability of labels displayed

```
with col1:
    st.markdown('Draw a letter here:')
    # Create a drawing canvas with desired properties
    canvas_result = st_canvas(
        fill_color="#ffffff",
        stroke_width=10,
        stroke_color='#ffffff',
        background_color="#000000",
        height=200,
        width=200,
        drawing_mode='freedraw',
        key="canvas",
    )

with col2:
    # Show that the resized image looks like
    st.markdown("What the model see's as input:")
    if canvas_result.image_data is not None:
        img = cv2.resize(canvas_result.image_data.astype('uint8'), (28, 28))
        img_rescaling = cv2.resize(img, (200, 200), interpolation=cv2.INTER_NEAREST)
        st.image(img_rescaling)

# Generate the prediction based on the users drawings
if st.button('Predict'):
    x_user_input = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    x_user_input = x_user_input.reshape(1, 28, 28, 1) / 255 # Reshape and normalize
    pred = model.predict(x_user_input)
    pred_label = labels_dict[pred.argmax()]
    st.header(f'Predicted Label: {pred_label}')

# Create a Plotly barchart of the predicted probabilities
fig = create_probability_fig(pred)
st.plotly_chart(fig, use_container_width = True)
```

Live Demo

Colab Demo:

Letter Prediction

Draw a letter here:




Predict

What the model see's as input:



Predicted Label: e



Thank you!

A dark blue diagonal gradient bar that starts from the bottom left corner and extends towards the top right corner, covering the lower half of the slide.